

PCA as Gradient Descent

Andrew Draganov

February 8, 2022

1 PCA objective function and gradient

Given a high-dimensional dataset $X \in \mathbb{R}^{N \times D}$, we aim to identify a gradient descent objective to find points $Y \in \mathbb{R}^{N \times d}$ with $d \ll D$ such that Y is the result of performing PCA on X .

We start with formalizing PCA as minimizing the function

$$f_{PCA}(X, Y) = \|L(E^X - E^Y)L\|_F^2 \quad (1)$$

where $E^X \in \mathbb{R}^{N \times N}$ and $E^Y \in \mathbb{R}^{N \times N}$ are the pairwise squared distance matrix for X and Y and L is the $N \times N$ centering matrix. Optimizing this function inherently amounts to minimizing the Frobenius norm between the two matrices. We show two separate ways to calculate the gradient of f_{PCA} with respect to the low-dimensional points Y .

1.1 Element-wise gradient calculations

To identify the effect of gradient descent on Y , we can deconstruct the matrices into a set of element-wise operations. We can re-arrange the term in the Frobenius norm:

$$\begin{aligned} L(E^X - E^Y)L &= LE^XL - LE^YL \\ &= E^X - \bar{X}^{\rightarrow} - \bar{X}^{\downarrow} - E^Y + \bar{Y}^{\rightarrow} + \bar{Y}^{\downarrow} \end{aligned}$$

where \bar{X}^{\rightarrow} is the $N \times N$ matrix of row-means that the centering matrix on the right-hand-side of X subtracts. The other \bar{X} and \bar{Y} variables are defined accordingly. We can now rearrange these to obtain:

$$L(E^X - E^Y)L = E^X - E^Y - \Lambda$$

for $\Lambda = -\bar{X}^{\rightarrow} - \bar{X}^{\downarrow} + \bar{Y}^{\rightarrow} + \bar{Y}^{\downarrow} \in \mathbb{R}^{N \times N}$

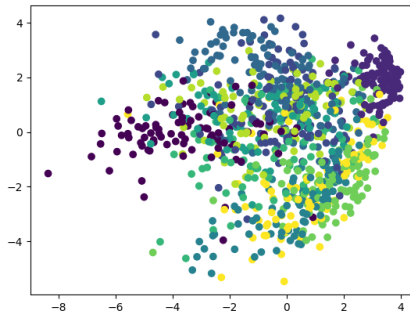


Figure 1: The result of performing gradient descent PCA on the MNIST dataset

Then the PCA objective function can equivalently be written as

$$f_{PCA}(X, Y) = \sum_{i,j} (||x_i - x_j||^2 - ||y_i - y_j||^2 - \lambda_{ij})^2 \quad (2)$$

To perform gradient descent, we take the partial derivative with respect to y_i to obtain

$$\nabla_{y_i} f(X, Y) = -4 \sum_j (||x_i - x_j||^2 - ||y_i - y_j||^2 - \lambda_{ij}) \left(\frac{\partial \lambda_{ij}}{\partial y_i} + \vec{y}_i - \vec{y}_j \right)$$

This gives us the following gradient descent algorithm with learning rate ν :

$$\begin{aligned} \nabla_{y_i} f(X, Y) &= -4 \sum_j (||x_i - x_j||^2 - ||y_i - y_j||^2 - \lambda_{ij}) \left(\frac{\partial \lambda_{ij}}{\partial y_i} + \vec{y}_i - \vec{y}_j \right) \\ Y_{t+1} &= Y_t + \nu \nabla_Y f(X, Y) \end{aligned}$$

This gives us the following algorithm:

Algorithm 1 PCA by gradient descent on the points

Require: Input: $X \in \mathbb{R}^{N \times D}$, n_epochs , ν

$D^X \leftarrow \text{pairwise_dists}(X) \in \mathbb{R}^{N \times N}$

$D^X \leftarrow D^X - \text{row_means}(D^X)$

$D^X \leftarrow D^X - \text{col_means}(D^X)$

$Y \leftarrow \mathcal{N}_{(0,1)} \in \mathbb{R}^{N \times d}$

while $e < n_epochs$ **do**

$V^Y \leftarrow \text{pairwise_vectors}(Y) \in \mathbb{R}^{N \times N \times 2}$

$D^Y \leftarrow \text{pairwise_dists}(Y) \in \mathbb{R}^{N \times N}$

$D^Y \leftarrow D^Y - \text{row_means}(D^Y)$

$D^Y \leftarrow D^Y - \text{col_means}(D^Y)$

$\nabla_Y \leftarrow -4 \cdot \text{SUM}((D^X - D^Y)V^Y, \text{axis} = 1)$

$Y \leftarrow Y + \nu \nabla_Y$

$e++$

end while

return Y

Lastly, we note that we can separate the PCA gradient descent algorithm into a pair of alternating forces. For $\Lambda^X = \bar{X}^{\rightarrow} + \bar{X}^{\downarrow}$ and $\Lambda^Y = -\bar{Y}^{\rightarrow} - \bar{Y}^{\downarrow}$ Namely, we have

$$\begin{aligned} \nabla_{y_i} f(X, Y) &= -4 \sum_{i,j} (||x_i - x_j||^2 - ||y_i - y_j||^2 - \lambda_{ij}) (\vec{y}_i - \vec{y}_j) \\ &= 4 \sum_{i,j} (||x_i - x_j||^2 - \lambda_{ij}^X) (\vec{y}_i - \vec{y}_j) - 4 \sum_{i,j} (||y_i - y_j||^2 - \lambda_{ij}^Y) (\vec{y}_i - \vec{y}_j) \\ &= 4 (F_X - F_Y) \end{aligned}$$

Each force is dominated by the large and small distances. However, we can deconstruct these into clear attractive and repulsive forces. If points x_i and x_j are one another's nearest neighbors in the high-dimensional space, then $||x_i - x_j||^2$ will always be less than its respective centering λ_{ij}^X . Thus, nearest neighbors in the high-dimensional space *only* exert attractions on the embedding. Similarly, farthest neighbors in the high-dimensional space *only* create repulsive forces. Inversely, nearest neighbors in the embedding exert repulsions while farthest neighbors exert attractions. PCA is the single convergence of these individual systems. Therefore, as one progresses from x_i 's nearest to its farthest neighbor, the force must this neighbor exerts on x_i must change monotonically.

1.2 Vectorized PCA gradient

A different approach would be to rely on matrix differentiation using vectorizations and Kronecker products. We first establish the following preliminaries:

1. Given Gram matrix $G_X = X^T X$, we can write E^X as $E^X = \text{diag}(G_X)\mathbb{1} + \mathbb{1}\text{diag}(G_X)^T - 2G_X$ as the matrix of pairwise squared distances

- We simplify notation by writing $\text{diag}(G_X) = g_x$, giving us

$$E^X = g_X \mathbb{1} + \mathbb{1}g_X^T - 2G_X$$

- $\mathbb{1}$ is the vector consisting of all 1's
2. The product of three matrices ABC can be *vectorized* (transformed into a column vector by linear transformation) using the operation $\text{vec}(ABC) = (C^T \otimes A)\text{vec}(B)$, where \otimes represents the Kronecker product.
 3. Vectorized matrices are transposed by appropriately sized *commutation matrices* K s.t. for any matrix $A \in \mathbb{R}^{n \times m}$ it is the case that $K^{(m \times n)}\text{vec}(A) = \text{vec}(A^T)$, where $K \in \mathbb{R}^{mn \times mn}$

We begin with the error function for PCA

$$f_{PCA}(X, Y) = \|L(E^X - E^Y)L\|_F^2$$

By applying the first preliminary, we rearrange this by

$$\begin{aligned} f_{PCA}(X, Y) &= \|L(E^X - E^Y)L\|_F^2 \\ f_{PCA}(X, Y) &= \|L(g_X \mathbb{1} + \mathbb{1}g_X^T - 2G_X - g_Y \mathbb{1} - \mathbb{1}g_Y^T + 2G_Y)L\|_F^2 \end{aligned}$$

But now notice that $L\mathbb{1} = 0 = \mathbb{1}^T L$. So we can simplify this expression into

$$f_{PCA}(X, Y) = \|L(2G_Y - 2G_X)L\|_F^2$$

The squared Frobenius norm $\|M\|_F^2$ can be rewritten as $\text{tr}(M^T M)$. Applying this gives us

$$\begin{aligned} f_{PCA}(X, Y) &= \text{tr} \left((L(2G_Y - 2G_X)L)^T L(2G_Y - 2G_X)L \right) \\ &= 4\text{tr} \left(L(G_Y - G_X)^T L L(G_Y - G_X)L \right) \\ &= 4\text{tr} \left(L(G_Y^T - G_X^T)L(G_Y - G_X)L \right) \\ &= 4\text{tr} \left(L(G_Y^T L G_Y - G_X^T L G_Y - G_Y^T L G_X + G_X^T L G_X)L \right) \end{aligned}$$

We now separate this by the linearity of the trace operation to obtain

$$f_{PCA}(X, Y) = 4 \left[\text{tr} (L G_Y^T L G_Y L) - \text{tr} (L G_X^T L G_Y L) - \text{tr} (L G_Y^T L G_X L) + \text{tr} (L G_X^T L G_X L) \right]$$

The product is commutative within the trace operation, so we can cancel the outer centering matrices to obtain

$$\begin{aligned} f_{PCA}(X, Y) &= 4 \left[\text{tr} (G_Y^T L G_Y) - \text{tr} (G_X^T L G_Y) - \text{tr} (G_Y^T L G_X) + \text{tr} (G_X^T L G_X) \right] \\ f_{PCA}(X, Y) &= 4\text{tr} (G_Y^T L G_Y - G_X^T L G_Y - G_Y^T L G_X + G_X^T L G_X) \\ &\quad \text{By further commutative swaps and removing the transposes, we get} \rightarrow \\ \rightarrow f_{PCA}(X, Y) &= 4\text{tr} (L (G_Y G_Y - 2G_X G_Y + G_X G_X)) \\ &= 4\text{tr} \left(L (G_Y - G_X)^2 \right) \end{aligned}$$

We now incorporate the fact that the trace of a product of symmetric matrices is the inner product of the vectorizations of those matrices

$$\rightarrow f_{PCA}(X, Y) = 4\langle \text{vec}((G_Y - G_X)^2), \text{vec}(L) \rangle$$

Taking the partial derivative with respect to Y gives:

$$\begin{aligned} d(f_{PCA}(X, Y)) &= 4\text{vec}(L)^T \text{vec} \left(d \left((G_Y - G_X)^2 \right) \right) \\ &= 4\text{vec}(L)^T \text{vec} \left(d \left(G_Y^2 \right) - d \left(G_Y G_X \right) - d \left(G_X G_Y \right) \right) \\ &= 4\text{vec}(L)^T \text{vec} \left(d \left(G_Y^2 \right) - 2G_X \right) \\ &= 4\text{vec}(L)^T \text{vec} \left(2G_Y d \left(G_Y \right) - 2G_X \right) \\ &= 8\text{vec}(L)^T \text{vec} \left(G_Y (dY^T Y + Y^T dY) - G_X \right) \\ &= 8\text{vec}(L)^T \text{vec} \left(G_Y (dY^T Y + Y^T dY) \right) - 8\text{vec}(L)^T \text{vec} (2G_X) \\ &= 8\text{vec}(L)^T \left(\text{vec} \left(G_Y dY^T Y \right) + \text{vec} \left(G_Y Y^T dY \right) \right) - 4\text{vec}(L)^T \text{vec} (2G_X) \end{aligned}$$

We now extract the dY vectors by performing the $\text{vec}(ABC) = (C^T \otimes A)\text{vec}(B)$ trick:

$$\begin{aligned} df_{PCA}(X, Y) &= 8\text{vec}(L)^T \left((Y^T \otimes G_Y)K dY + (I \otimes G_Y Y^T) dY \right) - 4\text{vec}(L)^T \text{vec} (2G_X) \\ &= 8\text{vec}(L)^T \left[((Y^T \otimes G_Y)K + (I \otimes G_Y Y^T)) dY \right] - 4\text{vec}(L)^T \text{vec} (2G_X) \end{aligned}$$